Eshita Kundu

Kolkata, India (open to remote) | eshita.kundu.2026@gmail.com
eshita.dev | github.com/eshitakundu | linkedin.com/in/eshitakundu

Summary

AI engineer (B.Tech CSE 2026, CGPA 8.8) building tool-using LLM agents where code, not the model, makes the final call. Shipped 3 deployed agents with adversarial evals in CI, 4 MCP servers and 8 merged open-source PRs (5 in AMD GAIA).

Technical Skills

LLMs & agents: Function calling, structured outputs, RAG, MCP / FastMCP, LangGraph, LangChain
Models & retrieval: OpenRouter, NVIDIA NIM, Claude and Gemini APIs, embeddings, FAISS vector search
Evaluation & testing: Agent evals in CI, adversarial / prompt-injection tests, grounding validators, pytest, Vitest, Playwright
Backend: Python, TypeScript, SQL, FastAPI, Pydantic, SQLAlchemy, REST APIs, React
Data & infrastructure: PostgreSQL, SQLite, Docker, GitHub Actions, Cloudflare Workers + D1, DigitalOcean, Linux

Projects

TraceIntel – Evidence-Grounded Security Agent Live · Demo · GitHub **2026**
Python, FastAPI, NVIDIA NOOA, OpenRouter, PostgreSQL, React, Docker, GitHub Actions

- Built an **LLM agent** that checks which token permissions a crypto transaction left open and advises revoke, verify or review, using **typed read-only tools** under a hard call budget.
- Grounded every claim: a validator rejects any sentence not backed by code-computed evidence, invalid output gets one repair attempt, and a **rule-based baseline** answers when the model fails.
- Wrote **24 CI evals** (**5 adversarial**: prompt injection, out-of-scope tool calls, fabricated claims) and **101 tests**; served by a FastAPI + PostgreSQL API with per-client rate limits, daily LLM cost caps and model-call traces.

ARIS – Multi-Branch Research Agent for Dependency Decisions Site · Demo · GitHub **2026**
Python, Heym, gpt-oss-120b on NVIDIA NIM, Tavily (MCP), Docker, DigitalOcean, Cloudflare

- Built a **104-node agent workflow** with **7 parallel research branches** over 8 sources; LLMs extract evidence and write an adopt / trial / hold / avoid brief.
- Kept numbers out of the model: **18 Python nodes** compute 6 weighted scores, a security veto and confidence.
- Built a **Python exec node for Heym** (7 tests, fork); proposed upstream, shipped by maintainers as a native Code node.

DevPulse – Grounded Opinion Agent over Hacker News Live · GitHub **2026**
Python, FastAPI, Pydantic, OpenRouter, Cloudflare Python Workers, D1, GitHub Actions

- LLM drafts pros and cons from HN threads; a **validator keeps only points quoted verbatim** from the cited comment.
- **CI guardrail eval** scripts 8 model failures (invented quotes, wrong citations, injection, timeouts, bad JSON): **0 ungrounded points published**; **21 tests**.

Agon – Guardrailed AI Trading Agent (team hackathon) Site · Demo **2026**
TypeScript, MCP, Solana, Jupiter API

- Co-designed an **MCP server** for AI clients to research tokens and paper-trade at live prices (**100+ accounts**); a deterministic guard outside the LLM enforces spend caps and slippage limits.

More: Study Buddy MCP (13 tools), GitHub Ranker remote MCP, RAG and multi-agent builds: agentic-ai list

Open Source

- **AMD GAIA** (local AI agents): **5 merged PRs, 1 open**. Made the **MCP HTTP client spec-compliant**, so official MCP SDK servers stop rejecting it (#6525); saved **MCP connectors** pick up catalog fixes (#6629); the **MCP bridge stops an LLM call when its client hangs up** (#6724); 2 docs fixes (#6238, #6239). Open: agents **pass large tool results between tools by handle**, outside the prompt (#8121).
- **CrossGL/crosstl** (shader compiler): **2 merged PRs** fixing WGSL codegen so brace initializers become typed constructors (#3672, #3692); **ai-engineering-from-scratch** (76k stars): #720 merged.

Experience

Employability.life | RPA Developer & Business Analyst (experiential learning) **Sep 2024 – Nov 2024**
Work-readiness program with Federation University Australia *Kolkata, India*

- Automated manual Excel reporting with UiPath, saving **10+ hours per week** and cutting errors by **70%**; wrote process design documents and led requirements, testing and sign-off across Agile sprints (certificate, recommendation).

Education

Sister Nivedita University – B.Tech, Computer Science and Engineering; CGPA: **8.80/10** **Jun 2022 – Jun 2026**

Achievements & Certifications

- **2nd Place, BRICS-FS-36 Data Analysis & Visualization**, international final, China, Dec 2024 (certificate)
- **EL Excellence Award, Top 50 Learner**, Federation University Australia, Feb 2026 (award, citation)
- **MCP Mastery**, Fractal Analytics / Coursera | **Applied Machine Learning in Python**, University of Michigan